

GraphSAGE Tutorial

Feb 18, 2021

Table of Contents

1. Installation
2. Guides
3. Data pre-processing
4. About the code
 - a. Train/ val/ test nodes and edges
 - b. Adjacency list
 - c. Identity features
 - d. Node sampling
 - e. Neighbor sampling
 - f. Test Stats
5. Execution script

Installation

Steps

1. git clone <https://github.com/williamleif/GraphSAGE>
2. (optional) create a virtual environment
 - a. pip install --user virtualenv
 - b. python -m venv [env_name]
3. pip install -r requirements.txt
 - a. If using python3:
 - b. Open requirements.txt
 - c. Replace futures==3.2.0 with futures==3.1.1

8 futures==3.2.0 → **8** futures==3.1.1

Guides


Guides

1. <https://github.com/williamleif/GraphSAGE> (README)
 - a. How to install
 - b. How to use
 - c. Models description
2. [Graphsage/example_data](#)
 - a. Example of the required data format
3. [Graphsage/example_supervised.sh](#) & [Graphsage/example_unsupervised.sh](#)
 - a. Example command to run the training and evaluation code in [graphsage/supervised_train.py](#)

Data

Karate Toy Data

- 34 nodes
- Node classification (4 classes)
- Node attributes (toy-data only):
 - Last name
 - Age
- Neighboring nodes:
 - High chance of same last name (e.g. siblings)
 - High change of similar age (± 5)



Preparing Data

Files: (data/prepare_data.py)

```
6 # Required:
7 # 1. Load the graph
8 # 2. Add validation and testing attributes to the graph
9 # 3. Create the features
10 # 4. Save the data in the proper format
```

```
35 # Prepare graph
36 nx.set_node_attributes(G, 'val', False)
37 nx.set_node_attributes(G, 'test', False)
38
39 val = ['32', '17', '6']
40 test = ['19', '10', '10', '6']
41
42 for node in val:
43     G.node[node]['val'] = True
44
45 for node in test:
46     G.node[node]['test'] = True
47
```

Generating Features

In this case, using:

1. Name
 - a. One-hot vector
2. Age
 - a. Integer
3. Node degree
 - a. Integer

```
48 # Generate features
49 feats = []
50
51 names_dict = {name: idx for idx, name in enumerate(set(names.values()))}
52 # print(names_dict)
53
54 for node in G:
55     name = names[node]
56     name_arr = [0]*len(names_dict)
57     name_arr[names_dict[name]] = 1
58     # 7-dim
59     # print(name_arr)
60
61     age = ages[node]
62     degree = G.degree()[node]
63
64     feat = name_arr + [age] + [degree]
65     # 9-dim
66     # print(feat)
67
68     feats.append(feat)
```

Writing Data

Write the processed data in the required format (json or npy).

All files should include the same prefix.

```
71 # Writing:
72 # Required files:
73 # 1. -G.json
74 # 2. -id_map.json
75 # 3. -class_map.json
76 # 4. -feats.npy
```

Writing Data

1. Graph written as json
2. ID map as json
 - a. It is required because nodes can be indexed by strings instead of integers.
3. Class map as json
4. Features as numpy array

```
81 # 1. Write graph
82 data = json_graph.node_link_data(G)
83 with open('processed/karate-G.json', 'w') as f:
84     json.dump(data, f)
85
86 # 2. Write id_map
87 id2idx = {n: idx for idx, n in enumerate(G.nodes())}
88 with open('processed/karate-id_map.json', 'w') as f:
89     json.dump(id2idx, f)
90
91 # 3. Write class_map
92 with open('processed/karate-class_map.json', 'w') as f:
93     json.dump(labels, f)
94
95 # 4. Write features
96 feats = np.array(feats)
97 # shape: [number of nodes, number of features]
98 np.save('processed/karate-feats.npy', feats)
```

GraphSAGE Code

Train/ Val/ Test Nodes and Edges

- Graph should have 'val' and 'test' attribute on every node.
- Edges connected to 'val' or 'test' nodes are **removed**.

```
45     broken_count = 0
46     for node in G.nodes():
47         if not 'val' in G.node[node] or not 'test' in G.node[node]:
48             G.remove_node(node)
49             broken_count += 1
```


```
55     for edge in G.edges():
56         if (G.node[edge[0]]['val'] or G.node[edge[1]]['val'] or
57             G.node[edge[0]]['test'] or G.node[edge[1]]['test']):
58             G[edge[0]][edge[1]]['train_removed'] = True
59     else:
60         G[edge[0]][edge[1]]['train_removed'] = False
61
```

Train/ Val/ Test Nodes and Edges

- Val and test nodes determined by the node attribute.
- The same node can be used for both val and test.
- Train nodes are all nodes except val and test.

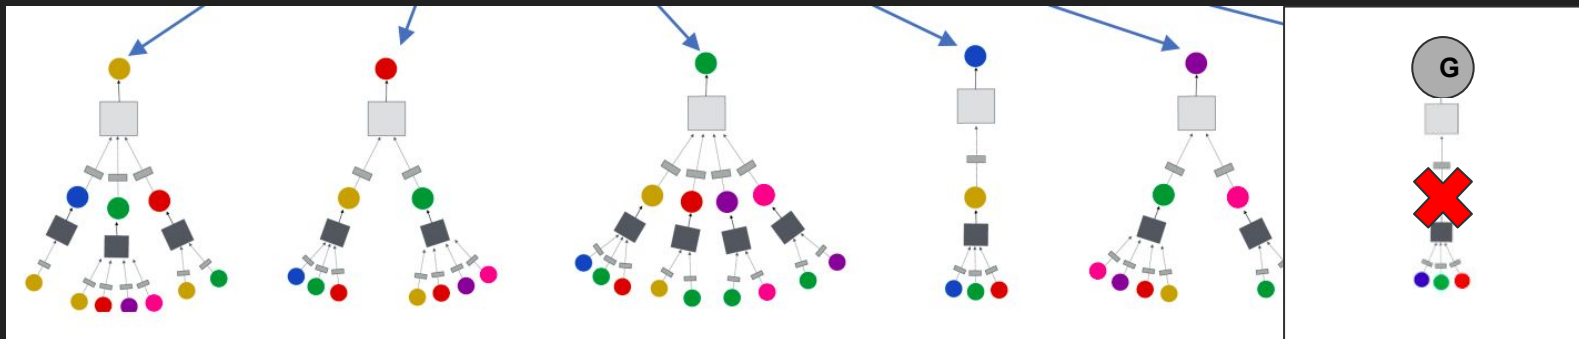
```
209     self.val_nodes = [n for n in self.G.nodes() if self.G.node[n]['val']]
210     self.test_nodes = [n for n in self.G.nodes() if self.G.node[n]['test']]
211
212     self.no_train_nodes_set = set(self.val_nodes + self.test_nodes)
213     self.train_nodes = set(G.nodes()).difference(self.no_train_nodes_set)
214     # don't train on nodes that only have edges to test set
215     self.train_nodes = [n for n in self.train_nodes if self.deg[id2idx[n]] > 0]
```

Implementation Challenge



Problem:
Isolated nodes have no
neighbors for aggregation

Solution:
Create a dummy vector to be the
neighbor for isolated nodes.



Adjacency List

Let max_degree = 128

Number of neighbors:

- If (>128):
 - Randomly choose 128 neighbors
- If (<128):
 - Randomly repeat some neighbors
- If (0):
 - Only points to dummy node

Files: minibatch.py, supervised_train.py

```
227 def construct_adj(self):
228     # Nodes: [0...n-1]
229     # n = len(id2idx)
230     adj = len(self.id2idx)*np.ones((len(self.id2idx)+1, self.max_degree))
231     deg = np.zeros((len(self.id2idx),))
232
233     for nodeid in self.G.nodes():
234         if self.G.node[nodeid]['test'] or self.G.node[nodeid]['val']:
235             continue
236         neighbors = np.array([self.id2idx[neighbor]
237                               for neighbor in self.G.neighbors(nodeid)
238                               if (not self.G[nodeid][neighbor]['train_removed'])])
239         deg[self.id2idx[nodeid]] = len(neighbors)
240         if len(neighbors) == 0:
241             continue
242         if len(neighbors) > self.max_degree:
243             neighbors = np.random.choice(neighbors, self.max_degree, replace=False)
244         elif len(neighbors) < self.max_degree:
245             neighbors = np.random.choice(neighbors, self.max_degree, replace=True)
246         adj[self.id2idx[nodeid], :] = neighbors
247     return adj, deg

133 if not features is None:
134     # pad with dummy zero vector
135     features = np.vstack([features, np.zeros((features.shape[1],))])
```

- Nodes #0...#(n-1)
 - Actual nodes
- Node #n
 - Dummy node


Identity Features

Files: supervised_models.py

- Identity features: randomly initialized features for each node.
- If “features” and “identity features” are both provided, they will be concatenated together.
 - E.g. 9-dimensional features + 50-dimensional identity features → 59-dimensional total features.
 - E.g. shape of “features” is (501, 9), shape of “embeds” is (501, 50) → shape of total features is (501, 59).

```
51     if identity_dim > 0:
52         self.embeds = tf.get_variable("node_embeddings", [adj.get_shape().as_list()[0], identity_dim])
53     else:
54         self.embeds = None
55     if features is None:
56         if identity_dim == 0:
57             raise Exception("Must have a positive value for identity feature dimension if no input features given.")
58         self.features = self.embeds
59     else:
60         self.features = tf.Variable(tf.constant(features, dtype=tf.float32), trainable=False)
61         if not self.embeds is None:
62             self.features = tf.concat([self.embeds, self.features], axis=1)
```

Node Sampling



Node Sampling


Files: minibatch.py

- (line 308) nodes are sampled from train_nodes in order.
- Every epoch, train_nodes is shuffled.

```
304     def next_minibatch_feed_dict(self):
305         start_idx = self.batch_num * self.batch_size
306         self.batch_num += 1
307         end_idx = min(start_idx + self.batch_size, len(self.train_nodes))
308         batch_nodes = self.train_nodes[start_idx : end_idx]
309         return self.batch_feed_dict(batch_nodes)

317     def shuffle(self):
318         """ Re-shuffle the training set.
319             Also reset the batch number.
320         """
321         self.train_nodes = np.random.permutation(self.train_nodes)
322         self.batch_num = 0
```

Neighbor Sampling



Neighbor Sampling

Files:

supervised_train.py

models.py

```
207     sampler = UniformNeighborSampler(adj_info)
208     layer_infos = [SAGEInfo("node", sampler, FLAGS.samples_1, FLAGS.dim_1),
209                   SAGEInfo("node", sampler, FLAGS.samples_2, FLAGS.dim_2)]
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254 def sample(self, inputs, layer_infos, batch_size=None):
255     """ Sample neighbors to be the supportive fields for multi-layer convolutions.
256
257     Args:
258         inputs: batch inputs
259         batch_size: the number of inputs (different for batch inputs and negative samples).
260     """
261
262     if batch_size is None:
263         batch_size = self.batch_size
264     samples = [inputs]
265     # size of convolution support at each layer per node
266     support_size = 1
267     support_sizes = [support_size]
268     for k in range(len(layer_infos)):
269         t = len(layer_infos) - k - 1
270         support_size *= layer_infos[t].num_samples
271         sampler = layer_infos[t].neigh_sampler
272         node = sampler((samples[k], layer_infos[t].num_samples))
273         samples.append(tf.reshape(node, [support_size * batch_size,]))
274         support_sizes.append(support_size)
275     return samples, support_sizes
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
```

Neighbor Sampling

Files: neigh_samplers.py

```
15 class UniformNeighborSampler(Layer):
16     """
17     Uniformly samples neighbors.
18     Assumes that adj lists are padded with random re-sampling
19     """
20     def __init__(self, adj_info, **kwargs):
21         super(UniformNeighborSampler, self).__init__(**kwargs)
22         self.adj_info = adj_info
23
24     def _call(self, inputs):
25         ids, num_samples = inputs
26         adj_lists = tf.nn.embedding_lookup(self.adj_info, ids)
27         adj_lists = tf.transpose(tf.random_shuffle(tf.transpose(adj_lists)))
28         adj_lists = tf.slice(adj_lists, [0,0], [-1, num_samples])
29         return adj_lists
```


Test Stats

File: supervised_train.py

```
326     print("Writing test set stats to file (don't peak!)")
327     val_cost, val_f1_mic, val_f1_mac, duration = incremental_evaluate(sess, model, minibatch, FLAGS.batch_size, test=True)
328     with open(log_dir() + "test_stats.txt", "w") as fp:
329         fp.write("loss={:.5f} f1_micro={:.5f} f1_macro={:.5f}".
330                 format(val_cost, val_f1_mic, val_f1_mac))
331
```

- Log dir is:
 - ./sup-[data folder name]
- Test results are saved there.

```
81 def log_dir():
82     log_dir = FLAGS.base_log_dir + "/sup-" + FLAGS.train_prefix.split("/)[-2]
83     log_dir += "{model:s}_{model_size:s}_{lr:0.4f}/".format(
84         model=FLAGS.model,
85         model_size=FLAGS.model_size,
86         lr=FLAGS.learning_rate)
87     if not os.path.exists(log_dir):
88         os.makedirs(log_dir)
89     return log_dir
```


Execution Script

Execution Script

In “GraphSAGE” folder:

- `python -m graphsage.supervised_train --train_prefix data/processed/karate`

Because the training data is small, adjust `batch_size`, `print_every`, etc.

- `python -m graphsage.supervised_train --train_prefix data/processed/karate --batch_size 4 --print_every 2 --validate_iter 5`

Done